

Final project presentation of Σ DAD

(Σ xtensible Dynamic Apache-Kafka Deployer)

By Group 15

What is Apache Kafka?



- Real time data streaming
- Website activity tracking
- Messaging and events listener
- Operational metrics
- Log aggregation
- Metrics

Most common use cases




Brij Kishore Pandey

DON'T FORGET TO SAVE

TOP 5 KAFKA USE CASES

Data Streaming



Log Aggregation




Message Queuing



Web Activity Tracker



Data Replication

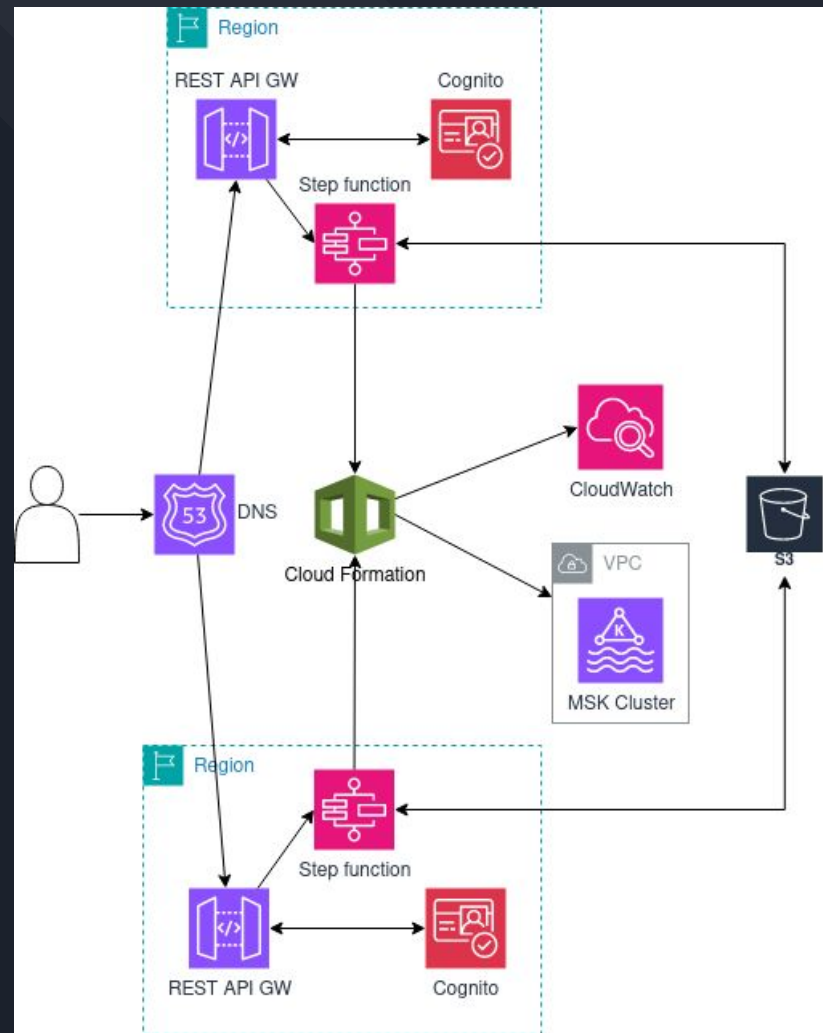




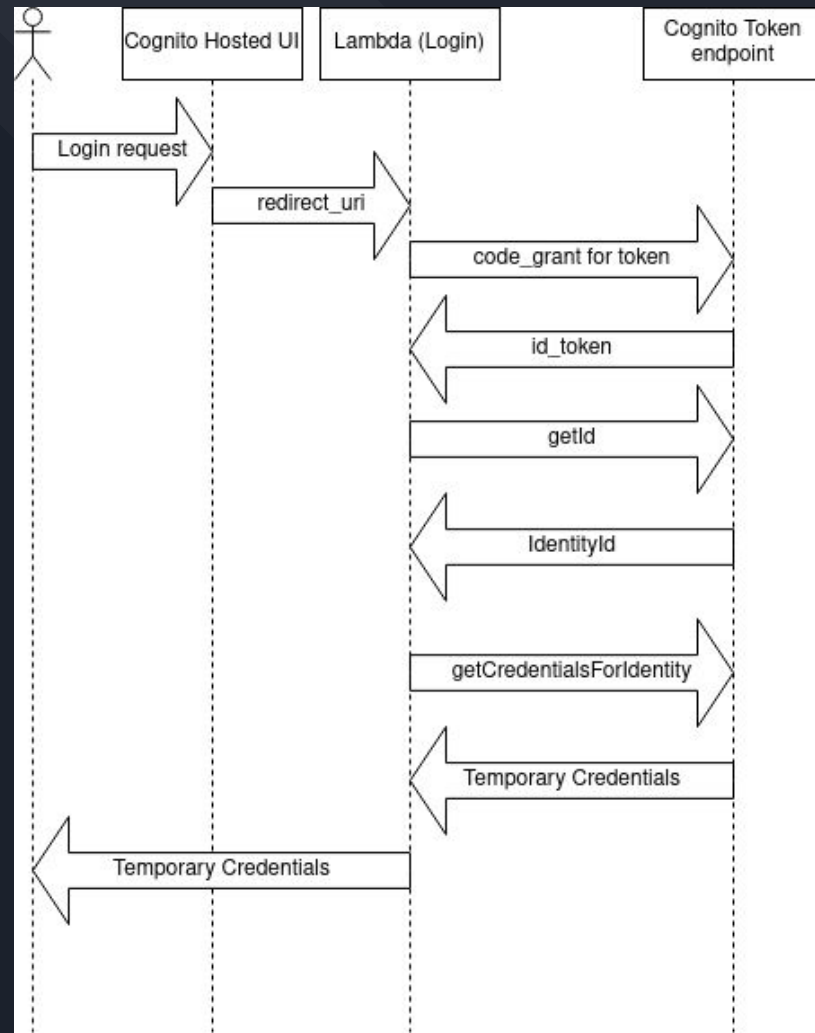
What problems we solve

- Removing tedious configuration of MSK clusters
- Learning new Cloud concepts
- Getting used to non-convenient UIs

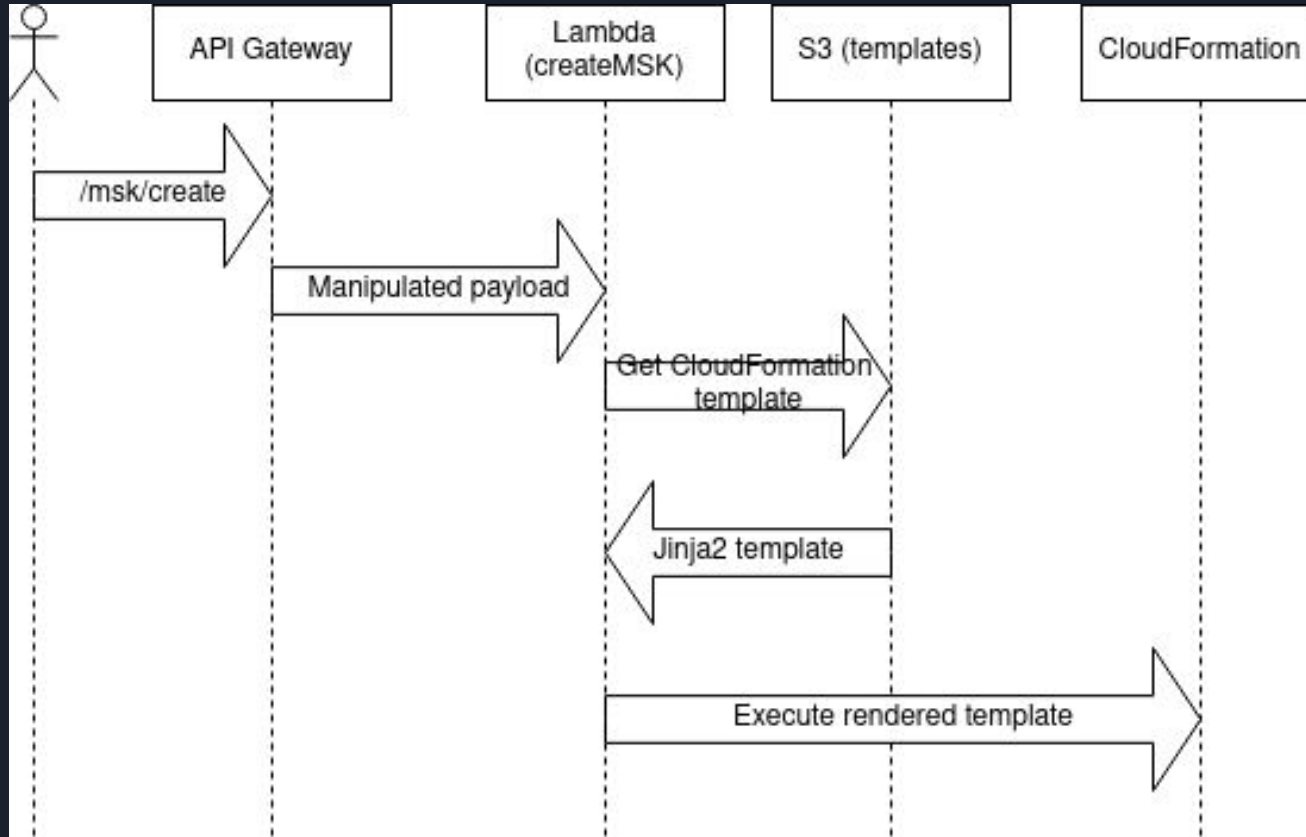
Network diagram




Authentication process for obtaining temporary credentials



Cluster creation Process



Cost awareness



Total cost

\$316.19


Average monthly cost

\$52.70

Service count

23

Costs (\$)



A more realistic Estimation

Estimate summary [Info](#)

Upfront cost
0.00 USD

Monthly cost
9,773.38 USD

Total 12 months cost
117,280.56 USD
Includes upfront cost

Getting Started with AWS

[Get started for free](#)[Contact Sales](#)

My Estimate

[<](#) **1** [>](#) [⊗](#)

<input type="checkbox"/>	Service Name	▼	Upfron... ▼	Monthly cost ▼	Descrip... ▼	Region ▼	Config Summary ▼
<input type="checkbox"/>	Amazon API Gateway		0.00 USD	972.36 USD	-	Europe (Fr...	REST API request units (exact number), Cache memory size (GB) (None), WebSocket message units (thousands), HTTP API reques...
<input type="checkbox"/>	Amazon API Gateway		0.00 USD	919.80 USD	-	Europe (Ir...	REST API request units (exact number), Cache memory size (GB) (None), WebSocket message units (thousands), HTTP API reques...
<input type="checkbox"/>	Amazon Cognito		0.00 USD	100.00 USD	-	Europe (Ir...	Advanced security features (Enabled), Number of monthly active users (MAU) (2000)
<input type="checkbox"/>	Amazon Cognito		0.00 USD	100.00 USD	-	Europe (Fr...	Advanced security features (Enabled), Number of monthly active users (MAU) (2000)
<input type="checkbox"/>	AWS Step Functions		0.00 USD	3,503.90 USD	-	Europe (Fr...	Workflow requests (6000 per hour), State transitions per workflow (32)
<input type="checkbox"/>	AWS Step Functions		0.00 USD	3,503.90 USD	-	Europe (Fr...	Workflow requests (6000 per hour), State transitions per workflow (32)
<input type="checkbox"/>	Amazon Managed Streaming for Apache Kafka (MSK)		0.00 USD	336.71 USD	-	Europe (Fr...	Storage per Broker (50 GB), DT Inbound: Internet (15 TB per month), DT Outbound: Amazon CloudFront (15 TB per month), DT I...
<input type="checkbox"/>	Amazon Managed Streaming for Apache Kafka (MSK)		0.00 USD	336.71 USD	-	Europe (Fr...	Storage per Broker (50 GB), DT Inbound: Internet (15 TB per month), DT Outbound: Amazon CloudFront (15 TB per month), DT I...



Security



AWS
Serverless



AWS IAM

The Cognito user Policy

```
{
  "Sid": "VisualEditor2",
  "Effect": "Allow",
  "Action": [
    "kafka-cluster:*Cluster",
    "kafka-cluster:Connect"
  ],
  "Resource": "arn:aws:kafka:eu-central-1:*:cluster/*/*",
  "Condition": {
    "StringEquals": {
      "aws:ResourceTag/MSK-Principal": "${aws:PrincipalTag/username}"
    }
  }
},
{
  "Sid": "VisualEditor3",
  "Effect": "Allow",
  "Action": [
    "kafka-cluster:*Data",
    "kafka-cluster:*Topic"
  ],
  "Resource": "arn:aws:kafka:eu-central-1:*:topic/*/*/*",
  "Condition": {
    "StringEquals": {
      "aws:ResourceTag/MSK-Principal": "${aws:PrincipalTag/username}"
    }
  }
},
{
```

```
{
  "Sid": "VisualEditor4",
  "Effect": "Allow",
  "Action": "kafka-cluster:*TransactionalId",
  "Resource": "arn:aws:kafka:eu-central-1:*:transactional-id/*/*/*",
  "Condition": {
    "StringEquals": {
      "aws:ResourceTag/MSK-Principal": "${aws:PrincipalTag/username}"
    }
  }
},
{
  "Sid": "VisualEditor5",
  "Effect": "Allow",
  "Action": [
    "kms:Decrypt",
    "kms:Encrypt",
    "kms:DescribeKey",
    "kms:Verify"
  ],
  "Resource": "*",
  "Condition": {
    "StringEquals": {
      "aws:ResourceTag/MSK-Principal": "${aws:PrincipalTag/username}"
    }
  }
},
{
```



Automation & Orchestration



AWS
CloudFormation



HashiCorp

Terraform



Some of the bigger Difficulties

- Rest API - VTL dynamic user Payload manipulation
- Cognito
- CloudWatch Dashboard sharing



Demo



Reflection

- Couldn't provide automatically Shareable Dashboard - deprecated API calls
- The brokers couldn't authenticate using IAM mechanism - tested with official approaches provided by AWS

Our growth



- Better understanding of the context under which different AWS services should be used
- What you definitely shouldn't do/use
- How to create/manage/interact with some of the most used services on AWS, including such out of the scope of the MOSCOW table
- We had many opportunities for bettering our troubleshooting skills - found inconsistencies and loopholes in the AWS Documentation

Thank you for your attention! And before asking anything, remember:

